

# Visualisation

David Barber

# What is Data Visualisation?

- Data is often very high dimensional meaning that we can't directly 'see' the data.
- It's difficult to get intuitions about the data since we can't 'see it'.
- In Data Visualisation we try to find a low dimensional representation (2 or 3 dimensions) so that we 'see something'.
- There is no 'correct' or 'perfect' visualisation. Every low dimensional representation will lose some information contained in the original high dimensional data.
- Such visualisations can be useful to see whether there might be clusters of datapoints, or which datapoints are in some sense 'similar' to another.
- Historically, methods such as PCA or Sammon 'mappings' were popular, but they are now less preferred.
- We tend to heuristically prefer representations that better preserve neighbourhood structure.
- Some visualisation methods (autoencoders) can be good but they are very expensive to train.

# Setup

Each data vector  $\mathbf{x}_n$  is in a high dimensional space. Given the set of datapoints

$$\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$$

we want to find a corresponding low dimensional (2 or 3) vector representation  $\mathbf{y}_n$  for each  $\mathbf{x}_n$  to give

$$\mathcal{Y} = \{\mathbf{y}_1, \dots, \mathbf{y}_N\}$$

- We would like  $\mathcal{Y}$  to preserve both the local and global structure in  $\mathcal{X}$ .
- Unfortunately, many 'classical' visualisation methods (Sammon mapping, Isomap, Locally Linear Embedding) don't work that well on real-world data sets.
- We will focus on Stochastic Neighbour Embedding (SNE) and its 'robust' variant t-SNE which is one of the most popular current approaches.

# Stochastic Neighbour Embedding (SNE)

We define an  $N \times N$  Markov transition matrix:

$$p_{j|i} = \frac{\exp\left(-(\mathbf{x}_i - \mathbf{x}_j)^2 / (2\sigma_i^2)\right)}{\sum_{j \neq i} \exp\left(-(\mathbf{x}_i - \mathbf{x}_j)^2 / (2\sigma_i^2)\right)}, \quad p_{i|i} = 0$$

Similarly we define

$$q_{j|i} = \frac{\exp\left(-(\mathbf{y}_i - \mathbf{y}_j)^2\right)}{\sum_{j \neq i} \exp\left(-(\mathbf{y}_i - \mathbf{y}_j)^2\right)}, \quad q_{i|i} = 0$$

- The transition  $p$  describes the neighbourhood structure – how easily can one jump to other points from a given point.
- If we want to preserve this structure, we need to find  $\mathcal{Y}$  such that  $q$  is approximately the same as  $p$ .
- Note that the  $\mathbf{y}_n$  scale is arbitrary, so we have ‘fixed’ the variance to  $1/\sqrt{2}$  in the  $y$  space.

# SNE

- SNE minimises the KL divergence between each conditional distribution  $p_i \equiv p_{\cdot|i}$ :

$$E(\mathcal{Y}) = \sum_i \text{KL}(p_i|q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$

- The optimisation is performed using gradient descent with parameters  $\mathcal{Y}$ .
- One can show (exercise) that

$$\frac{\partial E}{\partial \mathbf{y}_i} = 2 \sum_j (p_{j|i} - q_{j|i} + p_{i|j} - q_{i|j}) (\mathbf{y}_i - \mathbf{y}_j)$$

- Criticisms about this approach are that the KL divergence is not symmetric.
- For example there is a large cost for using widely separated  $y$  points (small  $q_{j|i}$ ) to represent nearby  $x$  points (large  $p_{j|i}$ ).
- SNE therefore focuses on making sure that the local structure is correct, but loses fidelity in retaining the global structure.
- Another problem is that the Gaussian form of  $p_{j|i}$  means that points which are far away will have negligible impact on the objective. We therefore need to make such points have an influence on the objective.

## t-SNE

There are many potential ways to 'robustify' the SNE procedure. The t-SNE approach, which seems to work reasonably well, is:

- t-SNE uses a 'symmetric' loss

$$E = \text{KL}(p|q) = \sum_i \sum_j p_{i,j} \log \frac{p_{i,j}}{q_{i,j}}$$

where the 'joint' distribution is defined

$$p_{i,j} = \frac{p_{j|i} + p_{i|j}}{2N}, \quad p_{i,i} = 0$$

- Note that this definition ensures that  $p_i = \sum_j p_{i,j} > 1/(2N)$  which encourages each datapoint to have a significant effect on the cost function.
- The gradient is (exercise)

$$\frac{\partial E}{\partial \mathbf{y}_i} = 4 \sum_j (p_{i,j} - q_{i,j}) (\mathbf{y}_i - \mathbf{y}_j)$$

- This works quite well, but there is an additional step that is also useful.

## t-SNE

- If we use a student t-distribution (rather than a Gaussian) this has heavier tails and can therefore assign non-negligible mass to  $y$  points that are quite far apart. Defining a student t-distribution with a single degree of freedom,

$$q_{i,j} = \frac{\left(1 + (\mathbf{y}_i - \mathbf{y}_j)^2\right)^{-1}}{\sum_{i \neq j} \left(1 + (\mathbf{y}_i - \mathbf{y}_j)^2\right)^{-1}}, \quad q_{i,i} = 0$$


Note that the sum in the denominator is over all pairs of distinct points.

- When  $(\mathbf{y}_i - \mathbf{y}_j)^2$  is large, the '1' term becomes irrelevant and the  $q$  distribution will be essentially invariant with respect to the overall length scale.
- Hence, for all but the finest length scales, pairs of points that are very far apart will have a similar contribution to points that are reasonably far apart.
- The gradient is (exercise)

$$\frac{\partial E}{\partial \mathbf{y}_i} = 4 \sum_j \frac{(p_{i,j} - q_{i,j})}{1 + (\mathbf{y}_i - \mathbf{y}_j)^2} (\mathbf{y}_i - \mathbf{y}_j)$$


- Note that is unclear why the t-SNE authors do not define a student t-distribution on the original  $x$  datapoints as well.

# MNIST 2D visualisation: t-SNE







# MNIST 2D visualisation: Sammon Mapping



## MNIST 2D visualisation: Isomap




## MNIST 2D visualisation: LLE



# Large Datasets

- Like most visualisation methods, t-SNE has an  $O(N^2)$  cost just to calculate the objective function since the distance between all pairs of points is calculated.
- For large datasets, this means that is very expensive to train t-SNE and related methods since each iteration of gradient descent requires an  $O(N^2)$  calculation.
- A cheaper alternative is to first define a desired number of neighbours and calculate a graph of which are the nearest neighbours of each node ( $x$  datapoint).
- We then select (randomly) a small set of 'landmark' datapoints in  $x$ , indexed by  $i'$ . We can then calculate a new transition matrix for these datapoints as follows. Starting from  $i'$  we randomly sample  $j$  according to  $p(j|i = i')$ . We repeat sampling from this Markov chain until we land on another landmark  $j' \neq i'$ . We then repeat this procedure many times for landmark  $i'$  and then normalise to obtain the transition  $p(j'|i)$ . We then repeat this for each landmark  $i'$ .
- We then use  $p(j'|i')$  in place of the original full  $p(j|i)$  transition to find a visualisation for the chosen landmark points.
- Whilst it is expensive to calculate  $p(j'|i')$ , this only needs to be done once.

## MNIST 2D visualisation: t-SNE



Visualisation of 6000 landmark points (using the random walk transition).

# Using Autoencoders to visualise MNIST

Autoencoders (with a 2D bottleneck) can also be used to visualise data).

**Fig. 3.** (A) The two-dimensional codes for 500 digits of each class produced by taking the first two principal components of all 60,000 training images. (B) The two-dimensional codes found by a 784-1000-500-250-2 autoencoder. For an alternative visualization, see (B).

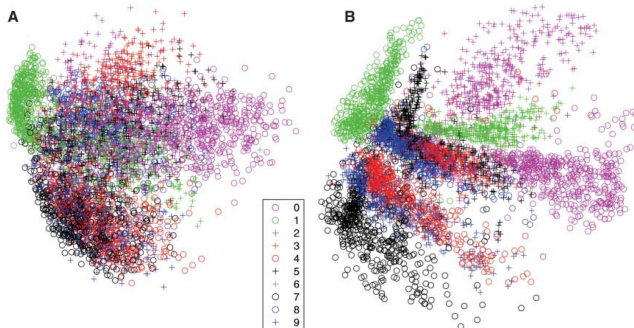


Figure from Hinton and Salakhutdinov Science 2006. The results can also be very good but training is quite slow and not practical for a 'quick' visualisation.