

STATG019 – Selected Topics in Statistics 2015

Lecture 1

An Introduction to Kernels

Dr Franz J. Király

Scope of the course

Part 1: Kernel Methods (Franz)

Hands-on crash course on kernels

Focus on „data analysis“ part

Selection of practical kernel methodology

Complements well topics in COMPGI13

(by Arthur Gretton, focus on learning theory and hypothesis testing)

Part 2: Spatial Point Processes (Simon)

Replace „kernels“ above with „spatial point processes“

Course organization

On-line course

moodle ID 16602

on-line activities, course info and materials

Lecture

Wednesdays, 9 - 11 am

Attendance is **mandatory**

First lecture Jan 14, last lecture Mar 25 (no lecture in reading week)

Slides and course script available on moodle

Lecture is videotaped and available on Lecturecast/moodle

Feel free to suggest content you want to learn about

Tutorials and/or practical sessions

Thursday, 11am - 1pm

location varies

May take place irregularly, depending on demand

Attendance is **non-compulsory**

Discussion of theoretical or practical exercises, group presentations, etc...

Your choice.

Table of contents

Lecture 1: Introduction to kernels

Main concepts and theoretical results, learning guarantees

Kernel PCA and kernel ridge regression

Some notes on R and kernlab

Lecture 2: the kernel support vector machine

The linear support vector machine, duality

Hard- and soft-margin two-class SVM

The one-class SVM

Support vector regression

Lecture 3: Gaussian processes and kernel learning

Potential further lecture topics:

Algorithms: kernel discriminants, kernel k-means, kernel quantile regression

kernel CCA, kernel MMD, kernel relevance vector machine

Large-scale learning with kernels, subset methods and Nyström-approximation

Combinatorial kernels: string kernels, graph kernels

Invariance kernels

Vapnik-Chervonenkis learning theory

Outlier detection, novelty detection


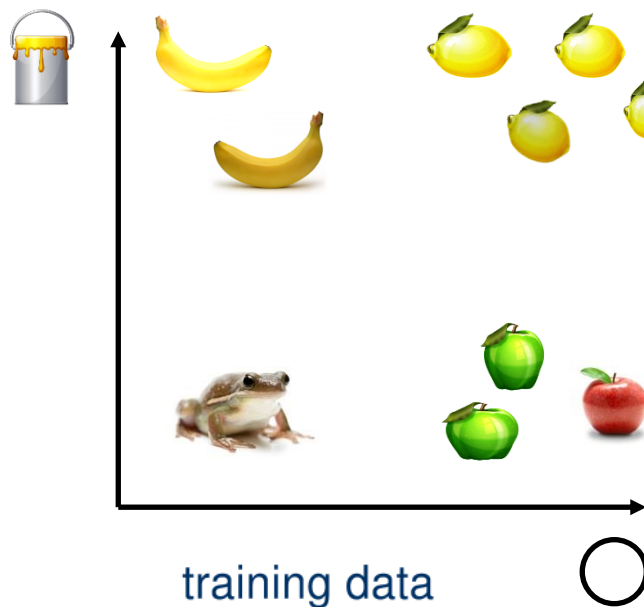
On-line kernel learning

An informal overview of things to come

Problem types in Statistical Machine Learning

Supervised

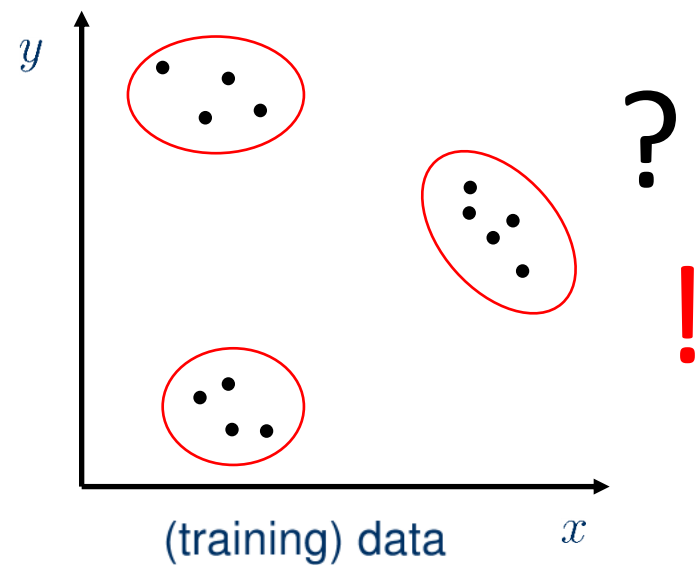

Some data is labelled by expert/oracle (mostly, training data)



Problem types in Statistical Machine Learning

Unsupervised

Data is not pre-labelled



What is Kernel Learning?

Most scalable algorithms employ linear principles

Most data exhibit non-linear features

Kernels allow to make linear algorithms work on non-linear features

Idea: replace scalar product in algorithm by “kernel product”

(this is neo-classical: Schölkopf 2002 – Learning with Kernels)

Example kernels:

$$k(x, y) = \langle x, y \rangle$$

“linear kernel” (nothing replaced)

$$k(x, y) = \exp \left(-\frac{1}{2\sigma^2} \|x - y\|^2 \right)$$

“Gauss kernel” (measures closeness)

$$k(x, y) = \langle x, y \rangle^d$$

“homogenous polynomial kernel

(measures shape)

$$k(x, y) = (\theta \langle x, y \rangle + 1)^d$$

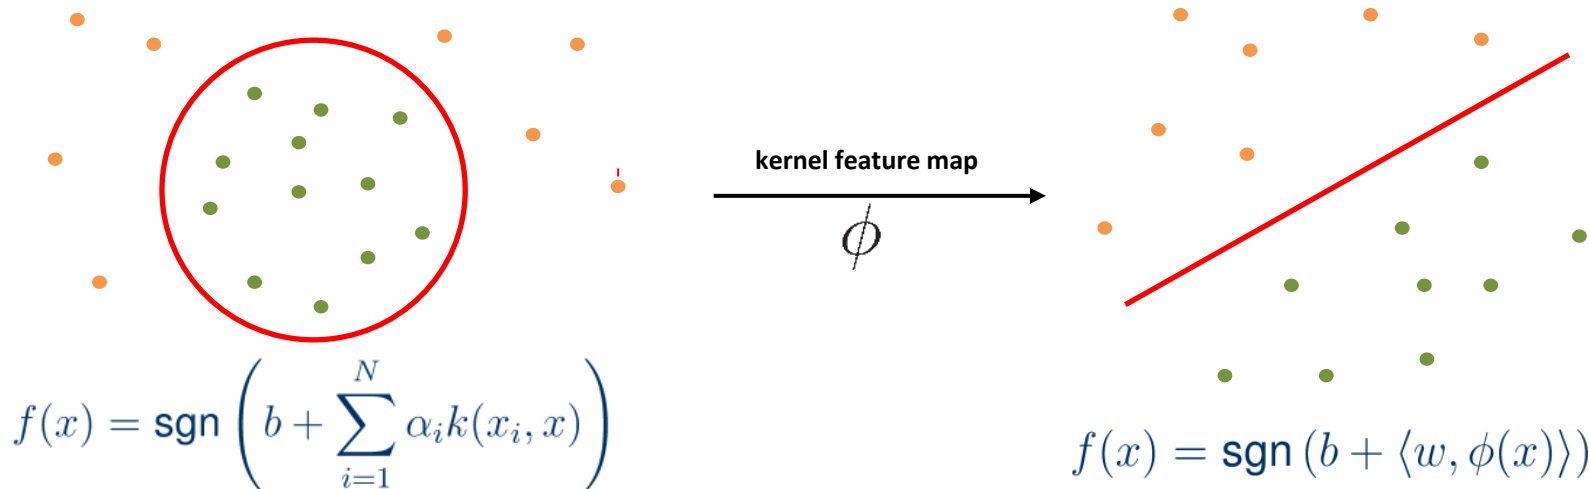
“inhomogenous polynomial kernel

The Kernel Trick

(classical, Aizerman et al 1964; Vapnik, 1995)

Kernel function is composition of “feature map” with scalar product

$$k(x, y) = \langle \phi(x), \phi(y) \rangle \quad \phi : \mathbb{R}^n \rightarrow \mathcal{F}$$



Trick = linearization of algorithms operating with scalar products
“kernelization”

Example: $k(x, y) = \langle x, y \rangle^2$ $\phi : (x_1, x_2) \mapsto (x_1^2, \sqrt{2}x_1x_2, x_2^2)$

“homogenous polynomial kernel” “Veronese map”

Kernelization

Kernels allow to make linear algorithms work on non-linear features

Idea: replace scalar product in algorithm by “kernel product”


Example: (ridge) regression

Input: $x_1, \dots, x_N \in \mathbb{R}^n, y_1, \dots, y_N \in \mathbb{R}$

Output: regressor function

$$f(x) = y^\top \cdot X (\lambda I + X^\top X)^{-1} \cdot x$$

$$X = (x_{i,j})_{ij} \in \mathbb{R}^{N \times n}$$



“Kernelized” variant:


$$f(x) = y^\top \cdot (\lambda I + K)^{-1} \cdot \kappa(x)$$

$$K = (k(x_i, x_j))_{ij}$$

“kernel matrix”

$$\kappa(x) = (k(x_i, x))_i$$

“kernel evaluation vector”



Kernelization

Example: support vector machine


Input: $x_1, \dots, x_N \in \mathbb{R}^n, y_1, \dots, y_N \in \{-1, +1\}$

Output: separator/decision function

$$f(x) = \text{sgn} (b + w^\top x)$$

$w = \sum_{i=1}^N \alpha_i y_i x_i$ solves a QP involving $X^\top X$

$\alpha_i, 1 \leq i \leq N$ are "dual" variables




“Kernelized”, non-linear variant:

$$f(x) = \text{sgn} \left(b + \sum_{i=1}^N \alpha_i k(x_i, x) \right) = \alpha^\top \kappa(x)$$

α solves a QP involving kernel matrix

$$K = (k(x_i, x_j))_{ij}$$



The kernel trick

Most scalable algorithms employ linear principles

Most data exhibit non-linear features

Naïve idea to make non-linear algorithms:

Suppose we are given data points $x_1, \dots, x_N \in \mathbb{R}^n$

Apply a non-linear “feature map” $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$ to all data points

Run your favourite *linear* algorithm on $\phi(x_1), \dots, \phi(x_N)$.

This idea (with a bit of care) works, and is at the basis of a considerable part of modern statistics!

E.g., non-linear Tikhonov-regularized least-squares regression:

$$f(x) = y^\top \cdot X \cdot (\lambda I + X^\top X)^{-1} \cdot x \quad X = \left(\phi(x_i)_j \right)_{ij} \in \mathbb{R}^{N \times m}$$

But there are issues with this approach:

Even for rather simple feature maps, m becomes quickly large.

In polynomial regression of degree d , one has $m = \Theta(n^d)$.

This causes: bad scaling/high runtime overfitting (even with regularization)
badly conditioned matrices

Also, the right choice of ϕ remains an issue.

Theorem (Singular Value Decomposition):

Every real matrix $M \in \mathbb{R}^{N \times n}$ admits a decomposition

$$M = U \cdot S \cdot V^T \quad \text{with} \quad \begin{aligned} U &\in \mathbb{R}^{N \times N} \text{ orthogonal} \\ V &\in \mathbb{R}^{n \times n} \text{ orthogonal} \\ S &\in \mathbb{R}^{N \times n} \text{ diagonal} \\ S &= \text{diag}(\sigma_1, \dots, \sigma_{\min(n, N)}) \\ \sigma_i &\text{ are non-negative} \\ &\text{(and usually ordered descendingly)} \end{aligned}$$

which is unique up to

orthogonal action on row-/column spaces
(if all σ_i are distinct, flipping sign in i -th column of U and V)

animation from Wikimedia Commons

columns of U, V : “left/right singular vectors”

σ_i : “singular values”

Corollary (by “uniqueness”):

eigenvectors of $M^T M$ = right singular vectors of M

left singular vectors of M = eigenvectors of $M M^T$


eigenvalues of $M^T M$ = (singular values)² of M = eigenvalues of $M M^T$

Inner product matrix!

This flavour of duality is at the core of “kernelization”.

Kernelizing principal component analysis

Principal Component Analysis: for data x_1, \dots, x_N



1. compute: sample mean $\hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i$
sample covariance matrix $\hat{\Sigma} = \frac{1}{N} \sum_{i=1}^N (x_i - \hat{\mu})(x_i - \hat{\mu})^\top$
2. $\text{eig}(\hat{\Sigma}_N)$ yields eigenpairs $(\lambda_1, v_1), \dots, (\lambda_n, v_n)$
 λ_i : “ i -th principal value” v_i : “ i -th principal vector”

“Principal scores”: $\tau_i(x) = v_i^\top (x - \hat{\mu}) / \sqrt{\lambda_i}$

Kernel Principal Component Analysis:

Observe: $\hat{\Sigma} = X_\mu^\top X_\mu$ where $X_\mu = (I - \mathbb{1}_N) X$ and

$$X = \begin{pmatrix} x_1 \\ \vdots \\ x_N \end{pmatrix} \in \mathbb{R}^{N \times n}$$

By SVD theorem, λ_i are the eigenvalues of

$\mathbb{1}_N$ the $(N \times N)$ matrix with entries $\frac{1}{N}$

$$K_\mu := X_\mu X_\mu^\top = (I - \mathbb{1}_N) \cdot X X^\top \cdot (I - \mathbb{1}_N) = (I - \mathbb{1}_N) \cdot K \cdot (I - \mathbb{1}_N)$$

where $K = (\langle x_i, x_j \rangle)_{ij}$ is the Gram matrix of the data.

easy kernelization via “kernel matrix” $K = (k(x_i, x_j))_{ij}$


By SVD theorem, v_i are right singular vectors of X_μ

So $\tau_i(x) = u_i^\top (X_\mu x - X_\mu X^\top \cdot \mathbb{1}/N) / \lambda_i$ where u_i is the corresponding left singular vector
 $= u_i^\top (\kappa(x) + (\mathbb{1}_N - I) \cdot K \cdot \mathbb{1}/N) / \lambda_i$ or equivalently, the corresponding eigenvector of K_μ
 where $\kappa(x) = (\langle x, x_i \rangle)$ is a cross-Gram vector and $\mathbb{1}$ is a vector of ones

easy kernelization via “empirical kernel vector” $\kappa(x) = (k(x, x_i))_i$

Kernelizing ridge regression

Ridge regression = Tikhonov-regularized OLS:



for data x_1, \dots, x_n $X = \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} \in \mathbb{R}^{N \times n}$

compute: $\hat{\beta} = (\gamma I + X^\top X)^{-1} X^\top \cdot y$

Prediction: $f(x) = \hat{\beta}^\top x$ minimizes $R(\beta) = \|y - X\beta\|_F^2 + \gamma \|\beta\|_2^2$
($\gamma \in \mathbb{R}_{>0}$ regularization parameter)

Kernel ridge regression:

Note: If (λ, v) is eigenpair of $A \in \mathbb{R}^{n \times n}$ then $(\lambda + \gamma, v)$ is eigenpair of $A + \gamma I$

Consider the singular value decomposition $USV^\top = X$ with $S = \text{diag}(\sigma_1, \sigma_2, \dots)$

with this, $(\gamma I + X^\top X)^{-1} X^\top = V \cdot S_\gamma \cdot U^\top$


where S_γ is diagonal with entries $\frac{\sigma_i}{\gamma + \sigma_i^2}$

By symmetry of the expression, $(\gamma I + XX^\top)^{-1} X = U \cdot S_\gamma \cdot V^\top$

So $\hat{\beta} = X^\top (\gamma I + XX^\top)^{-1} \cdot y = X^\top (\gamma I + K)^{-1} \cdot y$

$f(x) = x^\top \cdot X^\top \cdot (\gamma I + K)^{-1} \cdot y = \kappa(x)^\top \cdot (\gamma I + K)^{-1} \cdot y$

kernelization via $K = (k(x_i, x_j))_{ij}$ and $\kappa(x) = (k(x, x_i))_i$



As usual, $\hat{\beta} \rightarrow \hat{\beta}_{OLS}$ when $\gamma \rightarrow 0$ (Theorem)

if $m \geq N$, the limit $\gamma \rightarrow 0$ fits an *exact* regressor, i.e., $f(x_i) = y_i$

The Reproducing kernel Hilbert space formalism

The kernel trick in RKHS-generality

Previously: feature map $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$ $k(x, y) = \langle \phi(x), \phi(y) \rangle$
 positive definite kernel matrix $K = (k(x_i, x_j))_{ij}$ e.g. $k(x, y) = \langle x, y \rangle^d$

Observation 3: existence of ϕ has not been used in kernelization

All that is required is a well-behaved k that yields K

For example, $k(x, y) = \exp\left(-\frac{1}{2\sigma^2}\|x - y\|^2\right)$ always gives positive semi-definite K
 “Gauss kernel” and works very nicely for clustering.
 but there is no feature map $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$!

Definition: A function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ (usually $\mathcal{X} \subseteq \mathbb{R}^n$)

is called *positive definite kernel* if

(often required for technical reasons: $k \in L_2(\mathcal{X}^2)$ or $k \in L_\infty(\mathcal{X}^2)$)

every possible kernel matrix is symmetric and positive semi-definite.
 (in k and with data in \mathcal{X}) (sic)

Theorem (Moore-Aronszajn, 1950):

For every positive definite kernel k , there is (up to isomorphism)
 a unique feature map $\phi : \mathcal{X} \rightarrow \mathcal{H}$

into a unique Hilbert space \mathcal{H} , the so-called RKHS. (one way to imagine \mathcal{H} is $\mathbb{R}^{\mathbb{N}}$)

More precisely, $\phi(x) = (z \mapsto k(x, z))$ and $\langle k(x, \cdot), f(\cdot) \rangle_{\mathcal{H}} = f(x). \forall f \in \mathcal{H}$.

“reproducing (kernel Hilbert space) property”

A (non-exhaustive) list of popular kernels

$$k(x, y) = x^\top A y$$

linear kernels

$$k(x, y) = \rho(\|x - y\|)$$

Radial basis function kernels

$$k(x, y) = f(\langle x, y \rangle)$$

Dot-product kernels

$$k(x, y) = \langle x, y \rangle^d$$

homogenous polynomial kernel

$$k(x, y) = (\theta \langle x, y \rangle + 1)^d$$

inhomogenous polynomial kernel

(both polynomial kernels measure “shape”)

$$k(x, y) = \tanh(\sigma \langle x, y \rangle - \vartheta)$$

sigmoid kernel (measures “contours”)

not positive definite

$$k(x, y) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right)$$

Gaussian kernel

$$k(x, y) = \exp\left(-\frac{\|x - y\|}{\sigma}\right)$$

Laplacian kernel

$$k(x, y) = \left(1 + \frac{\|x - y\|^2}{\sigma^2}\right)^{-1}$$

Cauchy kernel

(these three measure “soft closeness”, with increasing long-range interaction/heavy tails from left to right)

Proposition:

For positive definite kernels k_1, k_2, \dots , and $\alpha_i \in \mathbb{R}_{\geq 0}$,

a k defined by $k(x, y) =$ as any below is a positive definite kernel:

$$k_3(x, y) + \alpha_{42}$$

$$\alpha_1 \cdot k_1(x, y) + \alpha_2 \cdot k_2(x, y)$$

$$k_1(x, y) \cdot k_2(x, y)$$

$$\sum_{\nu=1}^{\infty} \alpha_{\nu} \langle x, y \rangle^{\nu}$$

(in case of convergence)

$$\lim_{\nu \rightarrow \infty} k_{\nu}(x, y)$$

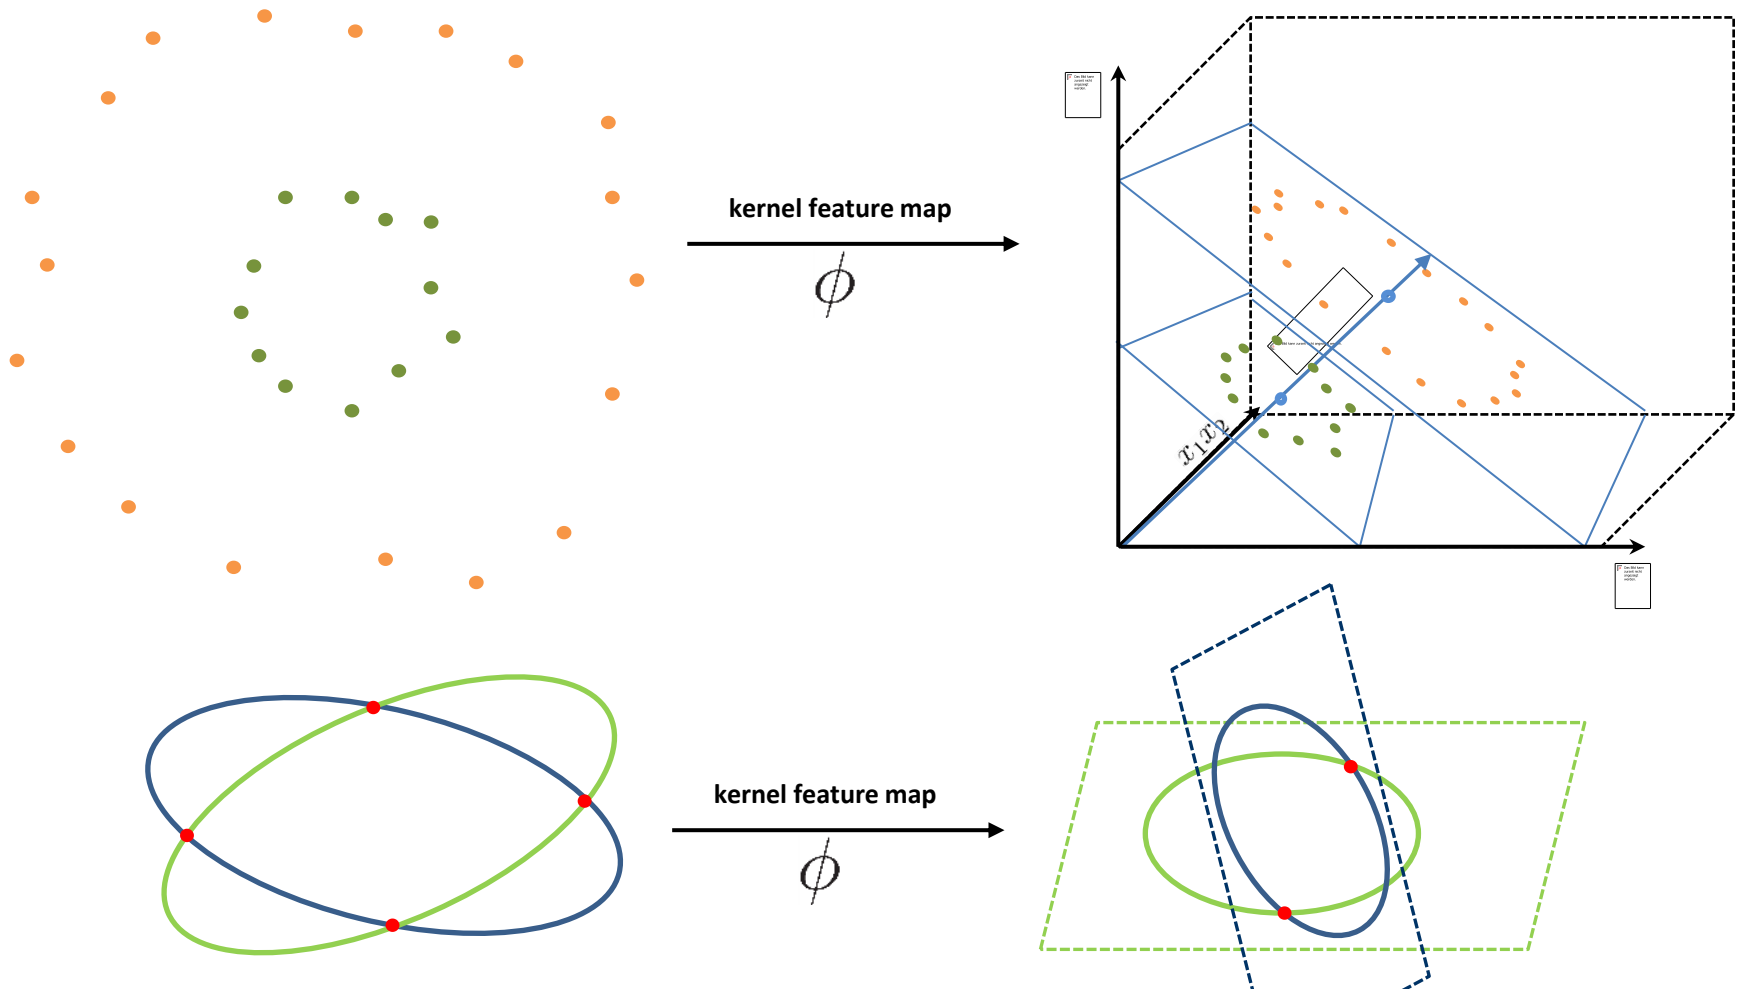
(in case of existence)

$$\frac{k_1(x, y)}{\sqrt{k_1(x, x) \cdot k_1(y, y)}}$$

Example: polynomial feature map

$$k(x, y) = \langle x, y \rangle^2 \quad \phi : (x_1, x_2) \mapsto (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$

$$k(x, y) = \langle \phi(x), \phi(y) \rangle \quad \phi : \mathbb{R}^n \rightarrow \mathcal{F}$$



Empirical risk minimisation and learning bounds

Recall: ridge regression finds a regressor function $f(x) = \hat{\beta}^\top x$
 where $\hat{\beta}$ minimizes $R(\beta) = \underbrace{\|y - X \cdot \beta\|_F^2}_{\text{"empirical risk"}} + \underbrace{\gamma \|\beta\|_2^2}_{\text{"regularizer"}}$

Similarly, kernel ridge regression finds a regressor $f(x) = \hat{\alpha}^\top \kappa(x)$
 where $\hat{\alpha}$ minimizes $R(\alpha) = \|y - K \cdot \alpha\|_{F, \mathcal{H}}^2 + \gamma \cdot \alpha^\top K \alpha$

One can show that we cannot do better under these circumstances:

Theorem (Kimeldorf-Wahba, 1971) “Representer Theorem”:

Let $(x_1, y_1), \dots, (x_N, y_N) \in \mathbb{R}^n \times \mathbb{R}$ be data points.

Let k be a positive definite kernel with RKHS \mathcal{H} .

Let $L : \mathbb{R}^n \times \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R} \cup \{\infty\}$ be any loss function,

for example squared loss $L(x_i, y_i, f(x_i)) = (y_i - f(x_i))^2$

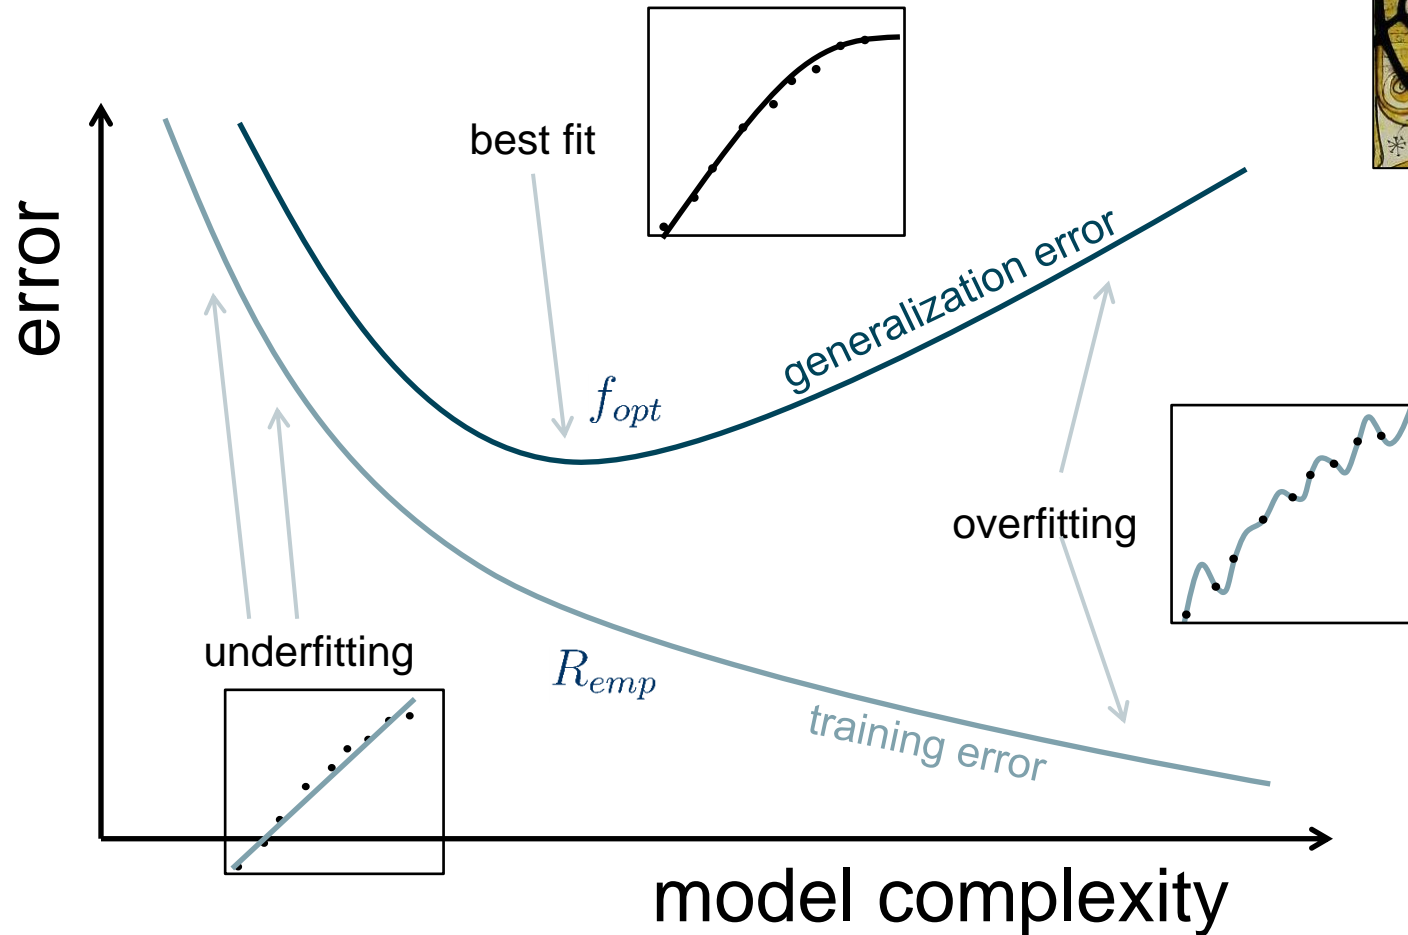
Let $\Omega : [0, \infty) \rightarrow \mathbb{R}$ strictly monotonously increasing. for example $\Omega(x) = \gamma \cdot x^2$

Then, there is a minimizer to $R(f) = \sum_{i=1}^N L(x_i, y_i, f(x_i)) + \Omega(\|f\|_{\mathcal{H}})$
 of the type $f(x) = \alpha^\top \kappa(x)$.

Statistical Learning Theory (Vapnik-Chervonenkis) quantifies

when and how $R_{emp}(f_{opt}) \rightarrow R(f_{opt})$ for $N \rightarrow \infty$.

VC = quantitative form of Occam's razor

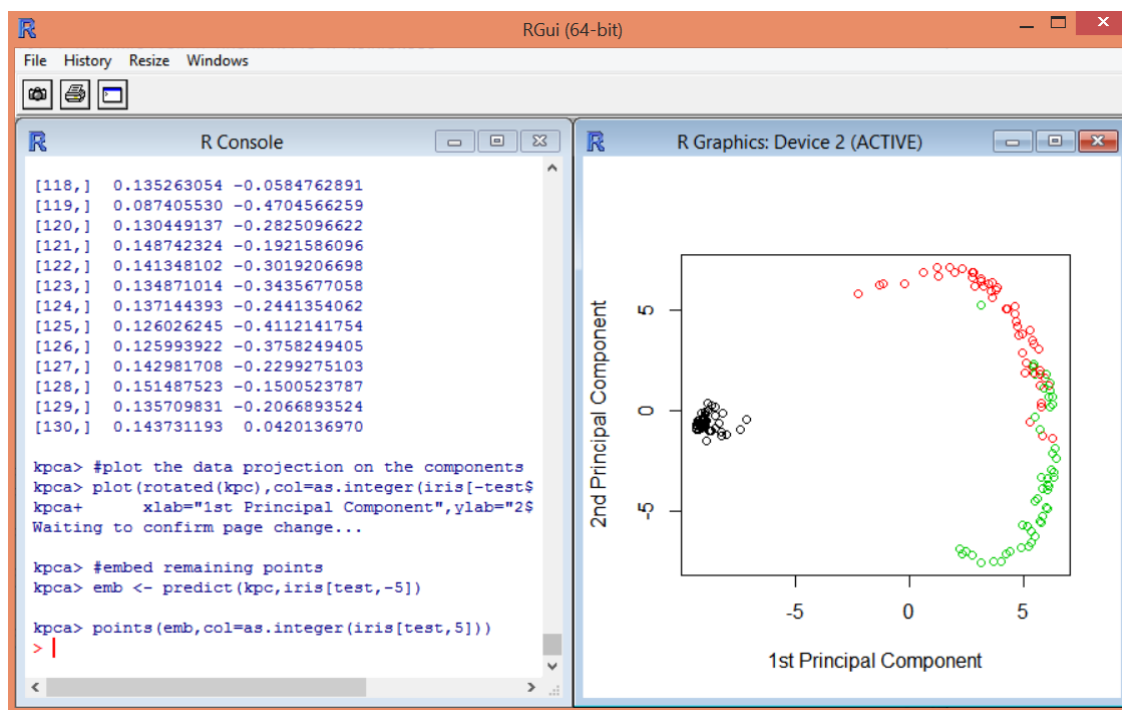


kernlab

Working with kernlab under R

If you have not worked with R before or not much:

look at STAT7001 materials or e.g. Dalgaard, introductory statistics with R
or type `help.start()` and click „an introduction to R“



type `install.packages(kernlab)` to install the kernlab package

After that `library(kernlab)` or `require(kernlab)` to use functions

Documentation: <http://cran.r-project.org/web/packages/kernlab/kernlab.pdf>

Try `help(functionname)` or `example(fuctionname)`

Outlook

Lecture 1: Introduction to kernels

Main concepts and theoretical results, learning guarantees

Kernel PCA and kernel ridge regression

Some notes on R and kernlab

Lecture 2: the kernel support vector machine

The linear support vector machine, duality

Hard- and soft-margin two-class SVM

The one-class SVM

Support vector regression

Lecture 3: Gaussian processes and kernel learning

Potential further lecture topics:

Algorithms: kernel discriminants, kernel k-means, kernel quantile regression

kernel CCA, kernel MMD, kernel relevance vector machine

Large-scale learning with kernels, subset methods and Nyström-approximation

Combinatorial kernels: string kernels, graph kernels

Invariance kernels

Vapnik-Chervonenkis learning theory

Outlier detection, novelty detection

On-line kernel learning

Next week: the Support Vector Machine

