

STATG019 – Selected Topics in Statistics 2015

Lecture 2

The Support Vector Machine

Dr Franz J. Király

Course organization

In-Course-Assessment

Two take-home ICA, one on kernels, one on point processes

Each counts 50% towards your final grade

Handing out: no.1 on Feb 9, no.2 on Mar 23 (on moodle)

Submission: no.1 on Mar 4, no.2 on Apr 29 (via moodle/Turnitin)

Further details will be announced via the moodle news forum

Tutorials and/or practical sessions

Thursday, 11am - 1pm

location varies

Given topic survey: next tutorial will be mainly on kernlab with R

Please install R and kernlab on your laptops (cluster rooms are full)

We do not have a room yet for tomorrow – do you know of any place?

Did you try the exercises? What do you think about them?

A short overview of the kernel SVM

the Support Vector Machine


Input: $x_1, \dots, x_N \in \mathbb{R}^n, y_1, \dots, y_N \in \{-1, +1\}$

Output: separator/decision function

$$f(x) = \text{sgn}(b + w^\top x)$$

$w = \sum_{i=1}^N \alpha_i y_i x_i$ solves a QP involving $X^\top X$

$\alpha_i, 1 \leq i \leq N$ are "dual" variables




“Kernelized”, non-linear variant:

$$f(x) = \text{sgn}\left(b + \sum_{i=1}^N \alpha_i k(x_i, x)\right) = \alpha^\top \kappa(x)$$

α solves a QP involving kernel matrix and vector

$$K = (k(x_i, x_j))_{ij} \quad \kappa(x) = (k(x_i, x))_i$$



The linear Support Vector Machine

(V. Vapnik, 1993)

The support vector machine

Goal: linearly separate two classes of points *with maximum margin*

Mathematically:

find a **hyperplane** $H = \{z \in \mathbb{R}^n : \langle a, z \rangle = c\}$

a the normal vector c the offset


separating the **green** and the **orange** points

with the largest symmetric margin of $\pm\eta$

$\langle a, x \rangle \geq c + \eta$ for all **orange** points x

$\langle a, x \rangle \leq c - \eta$ for all **green** points x

For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \{-1, 1\}$



This can be reformulated as the **optimization problem**

$$\max_{a, c, \eta} \eta \quad \text{s.t.} \quad y_i \cdot (\langle a, x_i \rangle - c) \geq \eta \quad \text{for all } 1 \leq i \leq N$$

Issue: overparamterization!

a, c, η can be multiplied by common factor

Solution: make substitutions $w := \frac{a}{\eta}$, $b := \frac{c}{\eta}$ thus $\eta = \frac{\|a\|}{\|w\|} = \frac{1}{\|w\|}$

to obtain $\max_{w, b} \frac{1}{\|w\|} \quad \text{s.t.} \quad y_i \cdot (\langle w, x_i \rangle - b) \geq 1 \quad \text{for all } 1 \leq i \leq N$

Issue: objective function inverse of square root
bad optimization behaviour

Solution: replace $\max_{w, b} \frac{1}{\|w\|}$ by the equivalent $\min_{w, b} \langle w, w \rangle = \|w\|^2$

The support vector machine

Goal: linearly separate two classes of points *with maximum margin*

Quadratic optimization program:

For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \{-1, 1\}$

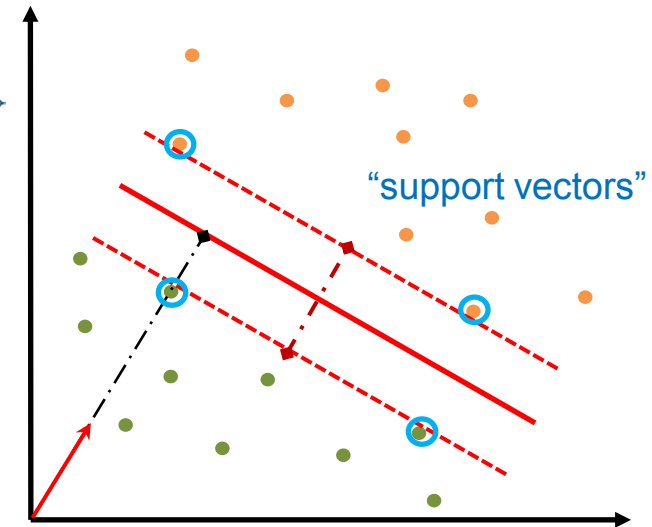
$$\min_{w, b} \frac{1}{2} \|w\|^2$$

$$\text{s.t. } y_i \cdot (\langle w, x_i \rangle - b) \geq 1 \quad \text{for all } 1 \leq i \leq N$$

$$\text{margin is } \pm \frac{1}{\|w\|}$$

$$\text{hyperplane is } H = \{z \in \mathbb{R}^n : \langle w, z \rangle = b\}$$

To an unseen data point $x \in \mathbb{R}^n$ assign the label $\text{sgn}(\langle w, x \rangle - b)$



“Naïve” Kernelization (Chapelle, 2006):

Replace x_i by $\phi(x_i)$ and w by $\sum_{i=1}^N \alpha_i \phi(x_i)$ (w is of this form by the representer theorem)

obtain:

$$\min_{\alpha} \frac{1}{2} \alpha^\top \cdot K \cdot \alpha$$

$$\text{s.t. } y_i \cdot (\alpha^\top \kappa(x_i) - b) \geq 1 \quad \text{for all } 1 \leq i \leq N$$

$$\text{where } k(y, z) = \langle \phi(y), \phi(z) \rangle$$

$$K = (k(x_i, x_j))_{ij} \quad \text{kernel matrix}$$

$$\kappa(x) = (k(x_i, x))_i \quad \text{empirical kernel vector}$$

This is a convex optimization problem in α

The following is about an alternative (“dual” SVM) which is more popular for historical reasons.

Lagrange duality

Lagrange duality

Start with an optimization program “*primal program*”

$$\min_{x \in \mathbb{R}^n} f(x)$$

where $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is the loss function to be minimized

$$\text{s.t. } g(x) \leq 0$$

(component-wise)

where $g : \mathbb{R}^n \rightarrow \mathbb{R}^N$ are N boundary conditions

x^* optimal solution (assumed to exist)

Definition: the *Lagrangian* for the above problem is

$$L(x, \alpha) := f(x) + \alpha^\top g(x) \quad L : \mathbb{R}^n \times \mathbb{R}^N \rightarrow \mathbb{R}$$

the *Lagrange dual function* is


$$W(\alpha) := \inf_{x \in \mathbb{R}^n} L(x, \alpha) \quad W : \mathbb{R}^N \rightarrow \mathbb{R}$$

the (Lagrange) dual program is

$$\max_{\alpha \in \mathbb{R}^N} W(\alpha)$$

s.t. $\alpha \geq 0$

α^* optimal solution (one can show: exists)



Proposition: it holds that $W(\alpha^*) \leq f(x^*)$ “duality gap”

Theorem (Slater): If the primal problem is *convex*, i.e., if f and g are, and there is $x \in \mathbb{R}^n$ such that $g(x) < 0$, then

$$W(\alpha^*) = L(x^*, \alpha^*) = f(x^*) \quad \text{“strong duality holds”}$$

Lagrange duality

primal program

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & f(x) \\ \text{s.t.} \quad & g(x) \leq 0 \end{aligned}$$

$$\begin{aligned} L(x, \alpha) &:= f(x) + \alpha^\top g(x) \\ W(\alpha) &:= \inf_{x \in \mathbb{R}^n} L(x, \alpha) \end{aligned}$$

dual program

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^N} \quad & W(\alpha) \\ \text{s.t.} \quad & \alpha \geq 0 \end{aligned}$$

Theorem (Slater): If the primal problem is *convex*, i.e., if f and g are, and there is $x \in \mathbb{R}^N$ such that $g(x) < 0$, then

$$W(\alpha^*) = L(x^*, \alpha^*) = f(x^*) \quad \text{“strong duality holds”}$$

Observation: $L(x, \alpha)$ is convex for fixed α

so hopefully x^* can be efficiently obtained as the minimum of $L(x, \alpha^*)$ in x

Example 1: linearly constrained linear program

primal program

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & \langle c, x \rangle \\ \text{s.t.} \quad & A \cdot x \leq b \\ & c \in \mathbb{R}^n, A \in \mathbb{R}^{N \times n} \end{aligned}$$

$$\begin{aligned} L(x, \alpha) &= c^\top x + \alpha^\top \cdot (Ax - b) \\ &= (c^\top + \alpha^\top A)x - \alpha^\top b \\ W(\alpha) &= \begin{cases} -\alpha^\top b, & c = A^\top \alpha \\ -\infty & \text{otherwise} \end{cases} \end{aligned}$$

dual program

$$\begin{aligned} \min_{\alpha \in \mathbb{R}^N} \quad & \langle b, \alpha \rangle \\ \text{s.t.} \quad & \alpha \geq 0 \\ & A^\top \alpha = c \end{aligned}$$

dual program is in standard form

Lagrange duality

primal program

$$\begin{aligned} \min_{x \in \mathbb{R}^n} \quad & f(x) \\ \text{s.t.} \quad & g(x) \leq 0 \end{aligned}$$

$$L(x, \alpha) := f(x) + \alpha^\top g(x)$$

$$W(\alpha) := \inf_{x \in \mathbb{R}^n} L(x, \alpha)$$

dual program

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^N} \quad & W(\alpha) \\ \text{s.t.} \quad & \alpha \geq 0 \end{aligned}$$

Example 2: the support vector machine

primal program

$$\min_{(w,b) \in \mathbb{R}^{n+1}} \frac{1}{2} \langle w, w \rangle \quad \text{s.t.} \quad y_i \cdot (\langle w, x_i \rangle - b) \geq 1 \quad \text{for all } 1 \leq i \leq N$$

note: this is a convex program

$$L((w, b), \alpha) = \frac{1}{2} \|w\|^2 + \sum_{i=1}^N [\alpha_i - \alpha_i y_i (\langle w, x_i \rangle - b)]$$

due to convexity, minimizer is zero of first derivatives:

$$\frac{\partial L}{\partial w} = w - \sum_{i=1}^N \alpha_i y_i x_i$$

$$\frac{\partial L}{\partial b} = \sum_{i=1}^N \alpha_i y_i$$

implies $w^* = \sum_{i=1}^N \alpha_i^* y_i x_i$
(compare representer theorem!)

therefore $W(\alpha) = L\left(\sum_{i=1}^N \alpha_i y_i x_i, b, \alpha\right) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \cdot x_i^\top x_j$ (if $\frac{\partial L}{\partial b} = 0$, othw. $-\infty$)

dual program $\max_{\alpha \in \mathbb{R}^N} W(\alpha) \quad \text{s.t.} \quad \alpha \geq 0, \quad 0 = \sum_{i=1}^N \alpha_i y_i$

Kernelizable!

The kernel

Support Vector Machine

The kernel support vector machine

The linear support vector machine

For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \{-1, 1\}$

Primal formulation:

$$\min_{w, b} \frac{1}{2} \|w\|^2 \quad \text{s.t.} \quad y_i \cdot (\langle w, x_i \rangle - b) \geq 1 \quad \text{for all } 1 \leq i \leq N$$


Dual formulation:

$$\max_{\alpha} W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \cdot x_i^\top x_j$$

$$\text{s.t. } \alpha \geq 0, \quad 0 = \sum_{i=1}^N \alpha_i y_i$$

$$\text{Obtain } w = \sum_{i=1}^N \alpha_i y_i x_i, \quad b = \frac{1}{N} \sum_{i=1}^N (w^\top x_i - y_i)$$

Predicted label is $y(x) := \text{sgn}(\langle w, x \rangle - b)$



The Kernel SVM

$$\max_{\alpha} W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \cdot k(x_i, x_j) \quad \text{s.t. } \alpha \geq 0, \quad 0 = \sum_{i=1}^N \alpha_i y_i$$

$$= \|\alpha\|_{\textcircled{1}}^2 - \frac{1}{2} \cdot \alpha^\top \tilde{K} \alpha$$

"sparse norm"

where $\tilde{K} = (y_i y_j k(x_i, x_j))_{ij}$ (this is a psd kernel matrix)

$$\text{Obtain } b = \frac{1}{N} \sum_{i=1}^N k(w, x_i) - y_i$$

Predict $\text{sgn}(\alpha^\top \tilde{\kappa}(x) - b)$ where $\tilde{\kappa}(x) = (y_i k(x_i, x))_i$

Intuition on the non-linear kernel SVM

For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \{-1, 1\}$

$$\tilde{K} = (y_i y_j k(x_i, x_j))_{ij}$$

$$\max_{\alpha} W(\alpha) = \|\alpha\|_1 - \frac{1}{2} \cdot \alpha^\top \tilde{K} \alpha \quad \text{s.t. } \alpha \geq 0, \quad 0 = \sum_{i=1}^N \alpha_i y_i$$


$$\tilde{\kappa}(x) = (y_i k(x_i, x))_i$$

Obtain $b = \frac{1}{N} \sum_{i=1}^N k(w, x_i) - y_i$ Predict $y(x) := \text{sgn}(\alpha^\top \tilde{\kappa}(x) - b)$

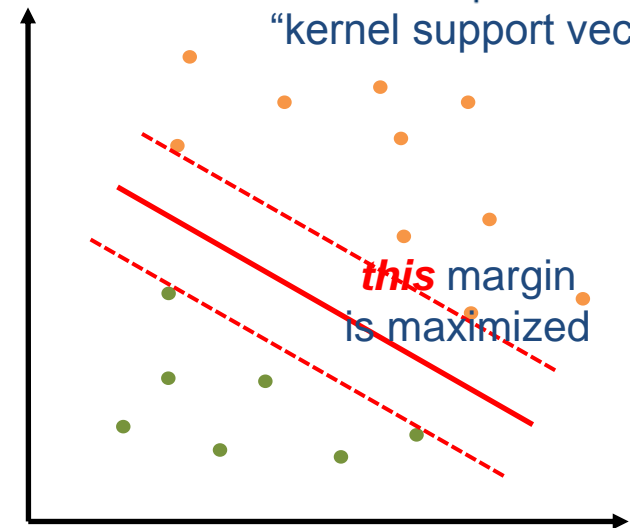
Observation: separator is manifold of form $\mathcal{M} = \{z : \alpha^\top \tilde{\kappa}(z) = b\}$

with α sparse

“kernel support vectors”



kernel feature map ϕ



For polynomial kernel: $\alpha^\top \tilde{\kappa}(z)$ is a polynomial $\sum_{i=1}^N \pm \alpha_i (\langle z, x_i \rangle)^d$ “sparse Waring form”

For Gaussian kernel: $\alpha^\top \tilde{\kappa}(z) = \sum_{i=1}^N \pm \alpha_i \exp\left(-\frac{\|z - x_i\|^2}{2\sigma^2}\right)$ due to feature space infinite:
 α often not very sparse
 (though many entries will be small)

A dictionary of kernel Support Vector Classifiers

name	how it differs from the vanilla-SVM	pros/cons	type in R kernlab-ksvm
(hard-margin) SVC	The original formulation (previous slide)	Very outlier sensitive, solution may not exist. Don't use.	C - svc set C=0
soft-margin SVC or C-SVC (Cortes, Vapnik)	$\min_{w,b,\xi} \frac{1}{2} \ w\ ^2 + C \sum_{i=1}^N \xi_i \quad \text{s.t.} \quad y_i \cdot (\langle w, x_i \rangle - b) \geq 1 - \xi_i$ $\xi_i \geq 0$ <p>data points can lie outside the margin but violation is punished infinitesimally by C</p>	Working horse of support vector classification. P: easy, sparse solution C: somewhat outlier sensitive	C - svc
nu-SVC (Schölkopf et al)	$\min_{w,\xi,\rho,b} \frac{1}{2} \ w\ ^2 - \nu \rho + \frac{1}{2} \sum_{i=1}^N \xi_i$ <p>parameter $\nu \in (0, 1)$</p> $\xi_i \geq 0$ $\text{s.t. } y_i \cdot (\langle w, x_i \rangle - b) \geq \rho - \xi_i \quad \rho \geq 0$ <p>data points can lie outside the margin fraction of violations of margin is hard-bounded above by ν</p>	support vector classification which emphasizes <i>number</i> of outliers above <i>severity</i> P: good for severe outliers C: tends to overfit more	nu - svc
bound-constraint SVC (Mangasarian et al)	$\min_{w,b,\xi} \frac{1}{2} \ w\ ^2 + \frac{1}{2} b^2 + C \sum_{i=1}^N \xi_i$ <p>otherwise same as C-SVC equivalent to homogenous & no intercept</p>	minor variant of C-SVC P: can be worth trying C: sometimes it isn't	C - bsvc
multi-class SVC	<p>Weston/Watkins: sum of C-SVC</p> $(\langle w, x_i \rangle - b)_{c(i)} \geq (\langle w, x_i \rangle - b + 2 - \xi_i)_{-c(i)}$ <p>Crammer/Singer: $\min_{M,\xi} \frac{\beta}{2} \ M\ _F^2 + \sum_{i=1}^N \xi_i$</p>	better than naïve one-vs-one or one-vs-all multiclassifiers W/W-risk is closer to C-SVC C/S-risk is closer to nu-SVC	kbb - svc spoc - svc

Regression and Outlier Detection with the SVM

Kernel support vector regression

Classification: For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \{-1, 1\}$
learn a classifier f such that $f(x_i) \approx y_i$


Regression: For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \mathbb{R}$
learn a regressor f such that $f(x_i) \approx y_i$

Idea: instead of mis-classification, penalize $\|f(x_i) - y_i\|$ with a *margin*

C-SVC:
$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \quad \text{s.t.} \quad y_i \cdot (\langle w, x_i \rangle - b) \geq 1 - \xi_i$$
$$\xi_i \geq 0$$

eps-SVR:
$$\min_{w,b,\xi^+,\xi^-} \frac{1}{2} \|w\|^2 + \frac{C}{N} \sum_{i=1}^N [\xi_i^+ + \xi_i^-]$$
$$\text{s.t.} \quad (\langle w, x_i \rangle - b) - y_i \leq \varepsilon + \xi_i^+, \quad \xi_i^+ \geq 0$$
$$y_i - (\langle w, x_i \rangle - b) \leq \varepsilon + \xi_i^-, \quad \xi_i^- \geq 0 \quad \text{for all } 1 \leq i \leq N$$

parameters
 $C, \varepsilon \in \mathbb{R}^+$



nu-SVR = eps-SVR + nu-SVC

$$\min_{w,b,\xi^+,\xi^-} \frac{1}{2} \|w\|^2 + C \left(\nu \varepsilon + \frac{1}{N} \sum_{i=1}^N [\xi_i^+ + \xi_i^-] \right)$$

parameters
 $C, \varepsilon \in \mathbb{R}^+$
 $\nu \in (0, 1)$

$$\text{s.t.} \quad (\langle w, x_i \rangle - b) - y_i \leq \varepsilon + \xi_i^+, \quad \xi_i^+ \geq 0$$
$$y_i - (\langle w, x_i \rangle - b) \leq \varepsilon + \xi_i^-, \quad \xi_i^- \geq 0$$

bound-constraint SVR
= eps-SVR + b^2

kernlab-ksvm:

eps-svr, nu-svr, eps-bsvr

(pros/cons analogous to SVCs)

Novelty detection with the one-class SVM

Classification: For data $x_1, \dots, x_N \in \mathbb{R}^n$ and labels $y_1, \dots, y_N \in \{-1, 1\}$
learn a classifier f such that $f(x_i) \approx y_i$

Novelty detection: Most data $x_1, \dots, x_N \in \mathbb{R}^n$ so far are “normal”
Learn a classifier that is able to identify “abnormal” points

Idea: “normal” points should be separated from origin by boundary (with slack)


C-SVC:
$$\min_{w, b, \xi} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \quad \text{s.t.} \quad y_i \cdot (\langle w, x_i \rangle - b) \geq 1 - \xi_i$$
$$\xi_i \geq 0$$

1-SVC:
$$\min_{w, \xi, \rho} \frac{1}{2} \|w\|^2 - \rho + \frac{1}{\nu \cdot N} \sum_{i=1}^N \xi_i$$
$$\text{s.t.} \quad \langle w, x_i \rangle \geq \rho - \xi_i, \quad \xi_i \geq 0 \quad \text{for all } 1 \leq i \leq N$$

parameter
 $\nu \in (0, 1)$

ν = upper bound on training points past the margin

Outlier score = distance to “normal” side $d(z) = \frac{w^\top x - \rho}{\|w\|}$



1-SVC after kernelization:

$$\min_{\alpha} \frac{1}{2} \alpha^\top K \alpha \quad \text{and} \quad d(z) = \frac{\alpha^\top \kappa(x) - \rho}{\sqrt{\alpha^\top K \alpha}} \quad \text{where} \quad K = (k(x_i, x_j))_{ij}$$
$$\text{s.t.} \quad 0 \leq \alpha \leq \frac{1}{\nu N}, \quad \|\alpha\|_1 = 1 \quad \kappa(x) = (k(x_i, x))_i$$

kernlab-ksvm: one-svc

improved variants exist (but not in R)

Using the SVM in kernlab

Support vector learning in kernlab

`ksvm` encapsulates all methods of support vector learning discussed today

Prior to use, kernlab has to be loaded with `library(kernlab)`

usage for training:

```
svmmodel <- ksvm(vartopredict~., data=traindata,...)
```

trains a SVM, stored in the output variable `svmmodel` of type `ksvm`

important parameters for `ksvm`:

<code>type</code>	determines the kind of learning machine, e.g. "C-svc"
<code>kernel</code>	determines the kernel used, e.g. "rbf-dot"
<code>kpar</code>	a list of kernel parameters, e.g. <code>list(sigma=1)</code>
<code>C, nu, epsilon</code>	regularization parameters for the various methods

output contains model as `alpha, b, alphaindex`:

usage for prediction:

```
predicted <- predict(svmmodel, testdata)
```

yields a vector `predicted` of predictions for `testdata`

Documentation: <http://cran.r-project.org/web/packages/kernlab/kernlab.pdf>

Type `help(ksvm)` or `example(ksvm)` for more details and examples

Methods for model selection


Major issue:

In kernels and in the support vector learning methods,
there are regularisation parameters to choose, such as

σ in Gaussian k. C in C-SVC ν in nu-SVC ε in SVR etc.

“universal” solution:

Cross-validation



for `ksvm`:

error estimation by k-fold cross-validation by setting param. `cross = k`
estimated cross-validation error is returned as `cross`

cross-validation for class probabilities with parameter `prob.model = T`
cross-validation for other parameters must be done manually (e.g. `cvTools`)

Outlook

Lecture 1: Introduction to kernels

- Main concepts and theoretical results, learning guarantees
- Kernel PCA and kernel ridge regression
- Some notes on R and kernlab

Lecture 2: the kernel support vector machine

- The linear support vector machine, duality
- Hard- and soft-margin two-class SVM
- The one-class SVM
- Support vector regression

Lecture 3: Gaussian processes and kernel learning

Potential further lecture topics:

- Algorithms: kernel discriminants, kernel k-means, kernel quantile regression
- kernel CCA, kernel MMD, kernel relevance vector machine
- Large-scale learning with kernels, subset methods and Nyström-approximation
- Combinatorial kernels: string kernels, graph kernels
- Invariance kernels
- Vapnik-Chervonenkis learning theory
- Outlier detection, novelty detection
- On-line kernel learning

Next week: Gaussian Processes

